

Mathematical Statistics Concepts

Data: X_1, \dots, X_n - result of some random experiment (lab experiment, survey sampling, ...)

Each X_i is random variable

Any transformation of the data $S = f(X_1, \dots, X_n)$ is called a statistic.

Examples: $S_1 = \min \{X_1, \dots, X_n\}$, $S_2 = \frac{1}{n} \sum_{i=1}^n X_i$ - are also random variables

Often we can assume that X_1, \dots, X_n are iid and have cdf $F(x) = P(X_1 \leq x)$. Given iid data X_1, \dots, X_n we want to learn something about F .

Example Lady tasting tea

Your friend claims that they can tell the difference between two similar drinks. You prepare n unlabeled cups, randomize the drinks and record $X_i \in \{0, 1\}$ - indicator of a successful drink identification at the i th trial.

$X_1, \dots, X_n \stackrel{\text{iid}}{\sim} \text{Bernoulli}(p)$,

where $P(X_i = 1) = p$, $P(X_i = 0) = 1 - p$. Let

Statistics $K = \sum_{i=1}^n X_i$ - # of successes

$\hat{p} = \frac{K}{n}$ is a reasonable estimator of p

$$E(\hat{p}) = E\left(\frac{K}{n}\right) = \frac{1}{n} E\left(\sum_{i=1}^n X_i\right) = \frac{1}{n} \sum_{i=1}^n E(X_i) = \frac{1}{n} \cdot n \cdot p = p$$

$$\text{Var}(X) = E[(X - E(X))^2]$$

$$\text{Var}(aX) = E[(aX - E(aX))^2] = E[a^2(X - E(X))^2] = a^2 E[(X - E(X))^2] = a^2 \text{Var}(X)$$

$$\text{Var}(\hat{p}) = \text{Var}\left(\frac{K}{n}\right) = \frac{1}{n^2} \text{Var}\left(\sum_{i=1}^n X_i\right) = [\text{independence}] = \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) = \frac{1}{n^2} \cdot n \cdot p(1-p) = \frac{p(1-p)}{n}$$



$$\text{Var}(\hat{p}) \xrightarrow[n \rightarrow \infty]{} 0 \quad \checkmark$$

Def Suppose $\hat{\theta}$ is an estimator of true quantity θ . $\text{Bias}(\hat{\theta}) \stackrel{\text{def}}{=} E(\hat{\theta} - \theta) = E(\hat{\theta}) - \theta$

If $E(\hat{\theta}) = \theta$, then $\hat{\theta}$ is called an unbiased estimator.

Monte Carlo Methods in Statistical Inference

When we observe data we summarize them into a set of statistics (e.g., $\hat{\theta} = \hat{\theta}(x_1, \dots, x_n)$). These statistics are random variables and we want to know their distributions. We ask ourselves "what if we repeated our experiment and obtained a new set of observations x_1^*, \dots, x_n^* , if we assume that we know cdf of X_1, \dots, X_n , we can "repeat" our experiment via Monte Carlo simulations.

Example MSE estimation

MSE = mean squared error

$$MSE(\hat{\theta}) = E[(\hat{\theta} - \theta)^2]$$

Pretending we know F - cdf of our data
we simulate m fake data sets

$$x_{11}^*, \dots, x_{1n}^* \rightarrow \hat{\theta}^{(1)}$$

\vdots

$$x_{m1}^*, \dots, x_{mn}^* \rightarrow \hat{\theta}^{(m)}$$

Monte Carlo estimate of MSE:

$$\widehat{MSE} = \frac{1}{m} \sum_{j=1}^m (\hat{\theta}^{(j)} - \theta)^2$$

Example Poisson rate estimation

$$X \sim \text{Poisson}(\lambda) \text{ if } \Pr(X=k) = \frac{\lambda^k}{k!} e^{-\lambda}, k=0,1,2,\dots$$

\uparrow
rate

$$E(X) = \lambda, \text{Var}(X) = \lambda.$$

Two estimators:

$$x_1, \dots, x_n \sim \text{Poisson}(\lambda)$$

$$\hat{\lambda}_1 = \bar{X} = \frac{1}{n} \sum_{i=1}^n x_i$$

\uparrow objective

$$\hat{\lambda}_2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{X})^2$$

Let's compare their MSEs

Monte Carlo Estimator of Confidence Interval Coverage.

Recall that $(1-\alpha)\%$ confidence interval (CI) promises to contain the true value of the parameter with probability $(1-\alpha)\%$.

We can check whether a CI lives up to its definition by Monte Carlo.

We generate m fake data sets and use them to form m CIs:

$$\begin{pmatrix} \hat{\theta}_l^{(1)} \\ \vdots \\ \hat{\theta}_l^{(m)} \end{pmatrix}, \begin{pmatrix} \hat{\theta}_u^{(1)} \\ \vdots \\ \hat{\theta}_u^{(m)} \end{pmatrix}$$

Coverage: $\widehat{\text{Coverage}} = \frac{1}{m} \sum_{j=1}^m \mathbb{1}_{\{\hat{\theta}_l^{(j)} < \theta < \hat{\theta}_u^{(j)}\}}$

$\widehat{\text{Coverage}}$ has a Monte Carlo error, so 0.95 ± 0.03 does not discredit our CI

Example: Poisson rate estimation

$$x_1, \dots, x_n \sim \text{Poisson}(\lambda)$$

$$\lambda_1 = \bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\hat{\sigma}^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2 \quad 95\% \text{ CI: } \hat{\lambda}_1 \pm 1.96 \frac{\hat{\sigma}}{\sqrt{n}}$$

How do we estimate Monte Carlo error of $\widehat{\text{Coverage}}$. Suppose $q = \text{Coverage}$ and $\hat{q} = \widehat{\text{Coverage}}$

$$\hat{q} = \frac{1}{m} \sum_{j=1}^m \mathbb{1}_h \quad \text{Monte Carlo error}$$

$$\sqrt{\frac{\hat{q}(1-\hat{q})}{m}}$$

$$\hat{q} \pm 1.96 \sqrt{\frac{\hat{q}(1-\hat{q})}{m}}$$

Monte Carlo for Hypothesis Testing

	H_0 true	H_0 false
reject H_0	Type I error α	Power $1-\beta$
fail to reject H_0	$1-\alpha$	Type II error β

We can estimate Type I error prob α or Power $1-\beta$ via Monte Carlo:

Type I error: simulate from H_0

Power: simulate from an alternative H_a
generate m data sets from H_0 :

$$x_{11}^*, \dots, x_{1n}^* \rightarrow T_1 \rightarrow I_1$$

$$x_{m1}^*, \dots, x_{mn}^* \rightarrow T_m \rightarrow I_m$$

test statistic

decision (reject H_0 or not)

$$\hat{\alpha} = \frac{1}{m} \sum_{j=1}^m I_j, \quad I_j = \begin{cases} 1 & \text{if rejected } H_0 \\ 0 & \text{if not} \end{cases}$$

if we change to H_a : $1-\hat{\beta} = \frac{1}{m} \sum_{j=1}^m I_j$

Bootstrap


Data: $x_1, \dots, x_n \stackrel{\text{iid}}{\sim} F \leftarrow \text{unknown}$

estimator: $\hat{\theta} = S(x_1, \dots, x_n)$

The main idea of bootstrap is to study the distribution of $\hat{\theta}^* = S(x_1^*, \dots, x_n^*)$, where

$x_1^*, \dots, x_n^* \sim \hat{F}$, where

\hat{F} is the empirical cdf defined as a cdf of a uniform random variable X taking values on $\{x_1, \dots, x_n\}$



$$\hat{F} \rightarrow F \text{ as } n \rightarrow \infty$$

If $X \sim \hat{F} \rightarrow$ simulating from \hat{F} is easy - sampling with replacement from $\{x_1, \dots, x_n\}$

Bootstrap estimation of estimator's bias

Given x_1, \dots, x_n and an estimator $\hat{\theta} = S(x_1, \dots, x_n)$ we would like to compute bias($\hat{\theta}$)

We construct B bootstrap samples by samp-

ling x_1, \dots, x_n with replacement n times

$$x_{11}^*, \dots, x_{1n}^* \leftarrow \text{sample 1}$$

\vdots

$$x_{B1}^*, \dots, x_{Bn}^* \leftarrow \text{sample } B$$

Notice that original x_i can appear more than once in each row above (i.e., we can have

$$x_{11}^* = x_3, \quad x_{15}^* = x_3)$$

Evaluate $\hat{\theta}^{(b)} = s(x_{b1}^*, \dots, x_{bn}^*)$ for $b = 1, \dots, B$

$$\overline{\hat{\theta}} \approx E(\hat{\theta}) = \frac{1}{B} \sum_{b=1}^B \hat{\theta}^{(b)}$$

$$\hat{\theta}_{\text{obs}} = s(x_1, \dots, x_n)$$

$$\widehat{\text{bias}}(\hat{\theta}) = \overline{\hat{\theta}} - \hat{\theta}_{\text{obs}}$$

Example: see lab quarto